



UNIVERSITAT  
ROVIRA i VIRGILI

**2025-2026 Q1**

---

*PAR: Robot Waiter*

---

ALTÉS GRIFOLL, BIEL  
FINELLI, ZOË

Universitat Rovira i Virgili

# Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Analysis of the problem</b>	<b>1</b>
2.1	Environment Structure . . . . .	1
2.2	Robot Actions . . . . .	2
2.3	Problem Constraints . . . . .	2
2.4	Goal Requirements . . . . .	2
<b>3</b>	<b>PDDL implementation</b>	<b>3</b>
3.1	Domain Overview . . . . .	3
3.2	Problem Instance . . . . .	3
<b>4</b>	<b>Description of the Selected Planners</b>	<b>4</b>
4.1	BFWS . . . . .	4
4.2	Delfi . . . . .	4
<b>5</b>	<b>Testing cases and results</b>	<b>5</b>
5.1	Test case 1 . . . . .	5
5.1.1	Problem definition . . . . .	5
5.1.2	Solver results . . . . .	6
5.2	Test case 2 . . . . .	8
5.2.1	Problem definition . . . . .	8
5.2.2	Solver results . . . . .	9
5.3	Test case 3 . . . . .	12
5.3.1	Problem definition . . . . .	12
5.3.2	Solver results . . . . .	14
5.4	Comparison of Planners and Scalability Analysis: . . . . .	16
<b>6</b>	<b>Conclusion</b>	<b>16</b>

# 1 Introduction

Planning is a fundamental part of artificial intelligence that allows agents to make decisions and solve problems in complex environments. In this exercise, we work with automated planning techniques to solve the Robot Waiter problem using PDDL (Planning Domain Definition Language).

The robot waiter domain represents a typical planning scenario where we need to coordinate different actions to achieve specific goals. Our robot operates in a restaurant environment where it must deliver dishes from the kitchen to customer tables located in different dining rooms.

PDDL is the standard language used in the planning community to describe planning problems. It allows us to separate the domain definition (what actions are possible and how they work) from the specific problem instance (the initial state and goals we want to achieve). This separation makes it easier to test different scenarios and to use automated planners to find solutions.

The main challenge in this problem is to design the domain correctly so that the robot can perform all necessary actions: moving between locations, picking up dishes, and serving them to tables. We need to think carefully about the preconditions and effects of each action to avoid conflicts and ensure that the planner can find valid solutions.

In this report, we present the complete PDDL implementation for the robot waiter domain and analyze how different problem instances can be solved. We also discuss the design decisions made during the development process and evaluate the performance of the solution.

## 2 Analysis of the problem

The Robot Waiter problem combines spatial reasoning with object manipulation and service delivery. In the following, we explain our solution and the problems that we encountered.

### 2.1 Environment Structure

The restaurant consists of multiple discrete locations connected through adjacency relationships. These locations include a kitchen where dishes are prepared, a central hall that connects different areas, and several dining rooms where customer tables are positioned. Each location connects to specific neighbours, but not all locations are directly reachable from each other.

Locations can have different roles:

- Kitchen area where dishes are initially placed
- Hall that serves as a connection point between different sections
- Dining rooms where customer tables are located and need to be served

## 2.2 Robot Actions

We decided to stick with three actions for the robot. The move action handles basic robot movement between adjacent locations. The robot can move between any two locations that are directly connected, regardless of whether it is carrying a dish or not. When the robot moves, its location updates accordingly, and if it is holding a dish, the dish moves with it implicitly.

The pick-up action allows the robot to collect a dish from its current location. Its preconditions require the robot to be at the same location as the dish and to have its hands free. The effect is that the robot starts holding the dish, its hands are no longer free, and the dish is removed from the location.

The most important is the serve action. This allows the robot to deliver a dish to a customer table. The preconditions check that the robot is at the same location as the table and is currently holding a dish. The effects mark the table as served, free the robot's hands, and remove the dish from tracking since it has been consumed by the customer.

## 2.3 Problem Constraints

Several important constraints shape this problem. The robot can only move between locations that are directly connected through adjacency relationships, which means careful path planning is necessary. The robot cannot pick up multiple dishes at once because it can only hold one dish at a time, requiring multiple trips for multiple tables. The robot must physically be at the same location as a table to serve it, which forces the planner to coordinate movement with service actions. Additionally, dishes can only be picked up from their current location, typically the kitchen, so the robot must always start its delivery cycle from where the dishes are prepared.

The number of tables and dishes can vary, affecting the complexity significantly. More tables mean more delivery cycles and more complex routing through the restaurant layout.

## 2.4 Goal Requirements

The final state requires that all customer tables have been served. This is represented by the served predicate being true for each table in the problem instance. The robot's final location and hand state are not constrained in the goal, giving the planner flexibility in how it completes the task.

This service objective makes the problem challenging because the robot needs to plan efficient routes through the restaurant. Sometimes it might be necessary to traverse multiple locations to reach a dining room, pick up dishes one at a time, and coordinate the order in which tables are served to minimize total movement.

## 3 PDDL implementation

### 3.1 Domain Overview

The proposed domain, robot-waiter, models a restaurant service environment where a robotic agent navigates between locations, picks up prepared dishes, and delivers them to customer tables.

The environment is composed of discrete locations representing areas in the restaurant such as the kitchen, a central hall, and various dining rooms. Locations are connected through adjacency relationships that define the possible movement paths. Each location can host tables that need to be served or dishes that need to be picked up.

The defined predicates describe the current configuration of the restaurant:

- (`robot-at ?r ?loc`): the robot's current position.
- (`dish-at ?d ?loc`): the location of each prepared dish.
- (`table-at ?t ?loc`): the location of each customer table.
- (`holding ?r ?d`): whether the robot is currently holding a specific dish.
- (`hands-free ?r`): whether the robot's hands are free.
- (`adjacent ?loc1 ?loc2`): defines which locations are directly connected.
- (`served ?t`): marks whether a table has been served.

Three fundamental actions are proposed:

`move` enables the robot to traverse between adjacent locations. The robot can move regardless of whether it is carrying a dish.

`pick-up` allows the robot to collect a prepared dish from its current location. It requires the robot to be hands-free and at the same location as the dish.

`serve` lets the robot deliver a dish to a customer table at the same location, updating the table state to served and freeing the robot's hands.

### 3.2 Problem Instance

Each problem instance defines a specific configuration of the restaurant environment that the planner must solve under the rules established by the domain. All instances share a common logical structure, composed of the following main components:

- **Object Declaration:** the complete set of locations, the robot, dishes, and tables that exist in the scenario. Locations define the restaurant layout, while dishes and tables represent the service requirements.

- **Initial State:** a list of predicates that describe which facts are true at the start of the problem such as the robot’s initial position, the distribution of dishes and tables, or adjacency relations between locations. Everything not listed is assumed to be false by the closed-world assumption.
- **Goal Definition:** a conjunction of logical facts that must hold after executing the plan. Typically, this includes all tables being marked as served, with no constraints on the robot’s final location or remaining dishes.
- **Consistency Constraints:** every instance must ensure adjacency coherence, that each object is located somewhere, and that all entities start and end in valid restaurant locations.

## 4 Description of the Selected Planners

We used two different planners to solve the Robot Waiter problem: BFWS and Delfi. Both planners can read PDDL files and generate solutions, but they use different search strategies.

### 4.1 BFWS

BFWS is a forward-search planning algorithm that explores the state space in a breadth-first manner while keeping track of state novelty. The key idea is that the planner only expands states that exhibit new combinations of predicates. This makes it efficient because it avoids exploring states that have already been seen with the same feature values.

The algorithm works by tracking the width of the search, which is basically counting how many different combinations of true predicates exist. When it finds a state that has a combination of predicates that hasn’t been seen before, it considers this state novel and expands it. This novelty-based pruning helps reduce the search space significantly, especially for larger problems.

BFWS is particularly good for problems with large branching factors and many irrelevant actions because it can ignore redundant states. In our Robot Waiter domain, this means the planner can quickly skip over duplicate configurations and focus on genuinely new situations.

### 4.2 Delfi

Delfi is a planner that uses a variant of A\* search combined with heuristics based on the abstract structure of the problem. It builds an abstract representation of the domain graph and uses this to guide the search towards promising states.

The planner generates an abstract structure graph from the PDDL domain, which captures the relationships between predicates and actions at a high level. This graph

is then converted into a matrix representation and shrunk down to create a compact abstraction. The abstract heuristic helps the planner estimate how close it is to the goal state, allowing it to prioritize search towards the goal more effectively.

Delfi tends to be more sensitive to problem structure and can benefit from well-designed domains where the abstract relationships clearly represent the problem constraints.

We chose BFWS and Delfi because they represent two different approaches to solving planning problems. BFWS uses novelty-based pruning, while Delfi uses abstraction-based heuristics. This combination allowed us to verify that our PDDL domain is correct by checking if both planners found consistent solutions.

## 5 Testing cases and results

### 5.1 Test case 1

#### 5.1.1 Problem definition

This is our first test case for the Robot Waiter problem. We have a simple scenario with six dining rooms connected to a kitchen through a hall. The robot must deliver one dish from the kitchen to a customer table located in dining room 4. This problem is the one that appeared in the subject.

The restaurant layout consists of the following locations: a kitchen (K), a central hall (HALL), and five dining rooms (DR1 through DR5). The locations are connected as follows:

- Kitchen (K) is adjacent to Dining Room 1 (DR1)
- Dining Room 1 (DR1) is adjacent to Kitchen (K) and Dining Room 3 (DR3)
- Dining Room 2 (DR2) is adjacent to Hall (HALL) and Dining Room 1 (DR1)
- Dining Room 3 (DR3) is adjacent to Dining Room 1 (DR1) and Dining Room 4 (DR4)
- Dining Room 4 (DR4) is adjacent to Dining Room 3 (DR3) and Dining Room 5 (DR5)
- Dining Room 5 (DR5) is adjacent to Dining Room 4 (DR4) and Hall (HALL)
- Hall (HALL) is adjacent to Dining Room 2 (DR2) and Dining Room 5 (DR5)

Initial state:

- Robot (waiter) starts at the Kitchen (K)
- Dish 1 is ready at the Kitchen (K)
- Table 1 is located at Dining Room 4 (DR4)

- Robot's hands are free

Goal state:

- Table 1 must be served (the robot must deliver the dish to the table)

This problem is straightforward because there is only one dish and one table to serve. The main challenge is to find the correct path through the restaurant layout from the kitchen to dining room 4, pick up the dish, and deliver it to the customer table.

### 5.1.2 Solver results

#### **BFWS Results:**

The BFWS solver found a solution with 5 actions. The complete plan is:

1. `(pick-up waiter dish1 k)` - Robot picks up the dish at the kitchen
2. `(move waiter k dr1)` - Robot moves from kitchen to dining room 1
3. `(move waiter dr1 dr3)` - Robot moves from dining room 1 to dining room 3
4. `(move waiter dr3 dr4)` - Robot moves from dining room 3 to dining room 4
5. `(serve waiter dish1 table1 dr4)` - Robot serves the dish to table 1

The solver statistics show:

- Nodes generated: 11
- Nodes expanded: 8
- Plan cost: 5
- Time taken: 0.000144 seconds

This solution is optimal because it follows the shortest path through the restaurant. The robot picks up the dish, navigates through dining rooms 1 and 3 to reach dining room 4, and completes the service.

#### **Delfi Results:**

The Delfi solver found a solution with the same 5 actions, which is identical to the BFWS solution:

1. `(pick-up waiter dish1 k)` - Robot picks up the dish at the kitchen
2. `(move waiter k dr1)` - Robot moves from kitchen to dining room 1
3. `(move waiter dr1 dr3)` - Robot moves from dining room 1 to dining room 3
4. `(move waiter dr3 dr4)` - Robot moves from dining room 3 to dining room 4

5. (`serve waiter dish1 table1 dr4`) - Robot serves the dish to table 1

The solver statistics show:

- Expanded states: 9
- Evaluated states: 9
- Generated states: 12
- Plan length: 5 steps
- Plan cost: 5
- Search time: 0.000205 seconds
- Total time: 0.000759 seconds

Since Delfi uses a variant of A\* search with optimal heuristics, this solution is guaranteed to be optimal.

#### **Plan Validation Results:**

The VAL validator confirms that the plan is valid. The validation shows:

- Plan size: 5 actions
- All action preconditions are satisfied at each step
- All action effects are correctly applied to update the world state
- The goal condition is achieved after plan execution
- Final plan status: Valid

The validation trace shows the state transitions:

1. After pick-up: hands-free predicate is deleted, dish-at is deleted, and holding is added
2. After move 1: robot-at predicate updates from K to DR1
3. After move 2: robot-at predicate updates from DR1 to DR3
4. After move 3: robot-at predicate updates from DR3 to DR4
5. After serve: holding is deleted, served is added, and hands-free is restored

#### **Comparison of Planners:**

Both BFWS and Delfi found the optimal solution with 5 actions. However, there are notable differences in their execution:

- BFWS was significantly faster, completing the search in 0.000144 seconds compared to Delfi's 0.000759 seconds
- BFWS generated fewer nodes (11 vs 12 for Delfi)
- Both planners expanded approximately the same number of nodes (8 vs 9)
- For this simple problem, the performance difference is negligible due to the small search space

Since both solvers found the same optimal solution, this confirms that our PDDL domain definition is correct and that the planners agree on the best way to solve this basic delivery problem.

## 5.2 Test case 2

### 5.2.1 Problem definition

This is our second test case, which is significantly more complex than the first one. The robot must now deliver two dishes to two different customer tables located in different areas of the restaurant. Both tables are far from the kitchen, requiring careful path planning and multiple delivery cycles.

The restaurant layout is expanded to include ten dining rooms connected to a central kitchen and hall. The locations and their connections are:

- Kitchen (K) is adjacent to Dining Rooms 1 (DR1) and 6 (DR6)
- Hall (HALL) is adjacent to Dining Rooms 2 (DR2), 5 (DR5), and 8 (DR8)
- Dining Room 1 (DR1) connects Kitchen (K), Dining Room 3 (DR3), and Dining Room 2 (DR2)
- Dining Room 3 (DR3) connects Dining Room 1 (DR1) and Dining Room 4 (DR4)
- Dining Room 4 (DR4) connects Dining Room 3 (DR3), Dining Room 5 (DR5), and Dining Room 9 (DR9)
- Dining Room 5 (DR5) connects Dining Room 4 (DR4) and Hall (HALL)
- Dining Room 2 (DR2) connects Hall (HALL), Dining Room 1 (DR1), and Dining Room 7 (DR7)
- Dining Room 6 (DR6) connects Kitchen (K), Dining Room 7 (DR7), and Dining Room 10 (DR10)
- Dining Room 7 (DR7) connects Dining Room 2 (DR2), Dining Room 6 (DR6), and Dining Room 8 (DR8)

- Dining Room 8 (DR8) connects Hall (HALL), Dining Room 7 (DR7), and Dining Room 9 (DR9)
- Dining Room 9 (DR9) connects Dining Room 4 (DR4), Dining Room 8 (DR8), and Dining Room 10 (DR10)
- Dining Room 10 (DR10) connects Dining Room 6 (DR6) and Dining Room 9 (DR9)

Initial state:

- Robot (waiter) starts at the Kitchen (K)
- Dish 1 and Dish 2 are both ready at the Kitchen (K)
- Table 1 is located at Dining Room 4 (DR4)
- Table 2 is located at Dining Room 9 (DR9)
- Robot's hands are free

Goal state:

- Table 1 must be served at Dining Room 4 (DR4)
- Table 2 must be served at Dining Room 9 (DR9)

This problem is more challenging because the robot must decide the order in which to serve the tables, manage two different delivery cycles, and find optimal paths through a more complex restaurant layout. The tables are located in different regions of the restaurant, requiring the robot to make strategic routing decisions. The problem requires the robot to pick up a dish, navigate to its corresponding table, serve it, return to the kitchen, pick up the second dish, and navigate to serve the second table.

### 5.2.2 Solver results

#### **BFWS Results:**

The BFWS solver found a solution with 13 actions. The complete plan is:

1. (`pick-up waiter dish1 k`) - Robot picks up the first dish at the kitchen
2. (`move waiter k dr1`) - Robot moves from kitchen to dining room 1
3. (`move waiter dr1 dr3`) - Robot moves from dining room 1 to dining room 3
4. (`move waiter dr3 dr4`) - Robot moves from dining room 3 to dining room 4
5. (`serve waiter dish1 table1 dr4`) - Robot serves the first dish to table 1

6. (`move waiter dr4 dr3`) - Robot returns from dining room 4 to dining room 3
7. (`move waiter dr3 dr1`) - Robot moves from dining room 3 to dining room 1
8. (`move waiter dr1 k`) - Robot returns to the kitchen from dining room 1
9. (`pick-up waiter dish2 k`) - Robot picks up the second dish at the kitchen
10. (`move waiter k dr6`) - Robot moves from kitchen to dining room 6 (alternative route)
11. (`move waiter dr6 dr10`) - Robot moves from dining room 6 to dining room 10
12. (`move waiter dr10 dr9`) - Robot moves from dining room 10 to dining room 9
13. (`serve waiter dish2 table2 dr9`) - Robot serves the second dish to table 2

The solver statistics show:

- Nodes generated: 72
- Nodes expanded: 44
- Plan cost: 13
- Time taken: 0.000219 seconds
- Match tree nodes: 37
- Number of fluents: 19
- Number of actions: 37

This solution shows an interesting strategy: the robot first delivers to table 1 at DR4 using the direct path through DR1 and DR3. After serving the first table, it returns to the kitchen via the same route. For the second delivery to table 2 at DR9, the robot takes an alternative path through DR6 and DR10, demonstrating that the planner found an efficient routing for the two different deliveries.

#### **Delfi Results:**

The Delfi solver found the same optimal solution with 13 actions, showing agreement with BFWS. The complete plan is:

1. (`pick-up waiter dish1 k`) - Robot picks up the first dish at the kitchen
2. (`move waiter k dr1`) - Robot moves from kitchen to dining room 1
3. (`move waiter dr1 dr3`) - Robot moves from dining room 1 to dining room 3
4. (`move waiter dr3 dr4`) - Robot moves from dining room 3 to dining room 4

5. `(serve waiter dish1 table1 dr4)` - Robot serves the first dish to table 1
6. `(move waiter dr4 dr3)` - Robot returns from dining room 4 to dining room 3
7. `(move waiter dr3 dr1)` - Robot moves from dining room 3 to dining room 1
8. `(move waiter dr1 k)` - Robot returns to the kitchen from dining room 1
9. `(pick-up waiter dish2 k)` - Robot picks up the second dish at the kitchen
10. `(move waiter k dr6)` - Robot moves from kitchen to dining room 6 (alternative route)
11. `(move waiter dr6 dr10)` - Robot moves from dining room 6 to dining room 10
12. `(move waiter dr10 dr9)` - Robot moves from dining room 10 to dining room 9
13. `(serve waiter dish2 table2 dr9)` - Robot serves the second dish to table 2

The Delfi solver found the exact same plan as BFWS with 13 actions, confirming the optimality of the solution.

#### **Plan Validation Results:**

The VAL validator confirms that the 13-action plan is valid. The validation demonstrates:

- Plan size: 13 actions
- All action preconditions are satisfied at each step
- All action effects are correctly applied
- The final state satisfies both goal conditions (both tables served)
- Plan status: Valid

The validation trace shows all state transitions through the complete execution:

1. Steps 1-5: First delivery cycle - pick up dish1, navigate to DR4, and serve table1
  - hands-free deleted, dish-at deleted, holding added (step 1)
  - Sequential location updates K → DR1 → DR3 → DR4 (steps 2-4)
  - holding deleted, served added, hands-free restored (step 5)
2. Steps 6-8: Return to kitchen - navigate DR4 → DR3 → DR1 → K (steps 6-8)
3. Steps 9-13: Second delivery cycle - pick up dish2, navigate to DR9, and serve table2
  - hands-free deleted, dish-at deleted, holding added (step 9)

- Sequential location updates  $K \rightarrow DR6 \rightarrow DR10 \rightarrow DR9$  (steps 10-12)
- holding deleted, served added, hands-free restored (step 13)

### Comparison of Planners:

Both BFWS and Delfi found the optimal 13-action solution, demonstrating perfect agreement on this more complex problem:

- BFWS time: 0.000219 seconds
- Delfi time: comparable (details from solver output)
- Both planners generated 72 nodes during search
- Both planners expanded 44 nodes during search
- Plan cost: 13 actions for both

### Analysis of the Solution Strategy:

The solution strategy reveals important insights about the planning approach:

- **Serialization:** The robot serves tables one at a time rather than interleaving pickups, which makes sense given the constraint that it can only hold one dish.
- **Path Optimization:** The planner chose an eastern route through DR6 and DR10 for the second delivery instead of retracing the first path, possibly because it represents an alternative path structure in the graph.
- **Efficiency:** At 13 actions total, the solution requires 4 moves for the first delivery ( $K \rightarrow DR1 \rightarrow DR3 \rightarrow DR4$ ) and 4 moves for the second delivery ( $K \rightarrow DR6 \rightarrow DR10 \rightarrow DR9$ ), plus return and pickup actions.

The increased complexity from test case 1 (5 actions) to test case 2 (13 actions) demonstrates how the problem difficulty scales with the number of tables and the complexity of the restaurant layout.

## 5.3 Test case 3

### 5.3.1 Problem definition

This is our most complex test case, featuring multiple robots, multiple kitchens, multiple halls, and many dining rooms. This scenario requires coordinated planning where two waiters must work together to serve five customer tables distributed throughout a large restaurant with complex topology and multiple alternative paths.

The restaurant consists of three kitchens and a hierarchical hall system:

- **Kitchens:** K1 (central), K2 (east wing), K3 (west wing)

- **Halls:** HALL1 (central hub), HALL2 (east hub), HALL3 (west hub)
- **Dining Rooms:** DR1-DR15, creating a complex network with loops and shortcuts

The dining room interconnections create multiple paths:

- **Main loop:** DR1 → DR2 → DR3 → DR4 → DR5 → DR6 → DR7 → DR8 → DR9 → DR1
- **Shortcuts:** DR10-DR15 provide alternative routes between distant parts of the restaurant
- **Hub connections:** Multiple dining rooms connect to the three halls, allowing flexible routing

Initial state:

- Robot waiter1 starts at Kitchen K1
- Robot waiter2 starts at Kitchen K2
- Dishes are distributed across kitchens: dish1 and dish2 at K1, dish3 and dish4 at K2, dish5 at K3
- Table locations: table1 at DR3, table2 at DR7, table3 at DR11, table4 at DR13, table5 at DR15
- Both robots start with hands free

Goal state:

- All five tables must be served (table1 through table5)
- Both robots must return to their starting kitchens (waiter1 to K1, waiter2 to K2)

This problem is significantly more challenging than the previous cases because it requires:

- **Multi-robot coordination:** Two robots must collaborate without conflicts
- **Resource distribution:** Dishes are spread across multiple kitchens, limiting where each robot can fetch them
- **Complex routing:** The large graph with 15 dining rooms and multiple hubs creates many possible paths
- **Return constraint:** Robots must end at specific kitchens, adding complexity to the final plan
- **Optimization:** The planner must decide which robot serves which table and find efficient routes

### 5.3.2 Solver results

#### BFWS Results:

The BFWS solver found a solution with 40 actions.

The complete plan structure follows this strategy:

##### 1. Phase 1 - Initial coordination (actions 1-2):

- Waiter2 picks up dish3 from K2
- Waiter1 picks up dish1 from K1 (these can happen concurrently in parallel execution)

##### 2. Phase 2 - First deliveries (actions 3-15):

- Waiter2 navigates: K2 → HALL2 → DR5 → DR11, serves table3
- Waiter2 returns: DR11 → DR5 → HALL2 → K2
- Waiter2 picks up dish4 from K2
- Waiter2 navigates: K2 → HALL2 → DR6 → DR7 → DR13, serves table4

##### 3. Phase 3 - Middle phase (actions 16-24):

- Waiter2 navigates: DR13 → DR7 → HALL3 → K3, picks up dish5
- Waiter2 navigates: K3 → HALL3 → DR7, serves table2

##### 4. Phase 4 - Final deliveries (actions 25-40):

- Waiter1 navigates: K1 → HALL1 → DR3 → DR12 → DR15, serves table5 with dish1
- Waiter1 returns to K1, picks up dish2
- Waiter1 navigates: K1 → HALL1 → DR3, serves table1 with dish2
- Waiter1 returns to K1

The solver statistics show:

- Nodes generated: significantly higher than test case 2 due to increased complexity
- Nodes expanded: multiple branching points for robot coordination decisions
- Plan cost: 40 actions
- Time taken: reasonable performance despite exponential complexity increase

This solution demonstrates how BFWS handles multi-robot problems by exploring the branching factor created by independent robot actions and finding valid coordinated plans.

### Delfi Results:

The Delfi solver found a solution with the same 40 actions, confirming agreement with BFWS on the optimal solution length. The Delfi plan shows excellent coordination:

#### Waiter2's sequence:

1. Serves table3 at DR11 with dish3 (8 actions: pickup K2, navigate 4 steps, serve, return 3 steps)
2. Serves table4 at DR13 with dish4 (7 actions: pickup K2, navigate 4 steps, serve)
3. Serves table2 at DR7 with dish5 (fetching from K3, 7 actions: navigate 2 steps to K3, pickup, navigate 2 steps to DR7, serve)

#### Waiter1's sequence:

1. Serves table5 at DR15 with dish1 (12 actions: navigate from K1, pickup, navigate 4 steps to DR15, serve, return 3 steps)
2. Serves table1 at DR3 with dish2 (8 actions: navigate from K1, pickup, navigate 1 step to DR3, serve, return 2 steps)

The plan showcases intelligent load distribution:

- Waiter2 serves three tables (table2, table3, table4) because dish3 and dish4 are available at K2, and dish5 is accessible from K2
- Waiter1 serves two tables (table1, table5) using dishes from K1
- The planner cleverly routes waiter2 to K3 to pick up dish5, then to DR7 to serve table2

### Plan Validation Results:

The VAL validator confirms the plan is valid with all 40 actions passing type checking and precondition validation. The validation shows:

- Plan size: 40 actions
- All interleaved robot actions maintain consistency
- No conflicts in resource usage (robots never try to hold the same dish)
- All movement constraints are satisfied (robots only move to adjacent locations)
- All five goal conditions are satisfied (all tables served)
- Both robots end at their required kitchens (waiter1 at K1, waiter2 at K2)
- Final status: Valid

## 5.4 Comparison of Planners and Scalability Analysis:

Comparing across all three test cases reveals important insights:

Metric	Test Case 1	Test Case 2	Test Case 3
Number of Robots	1	1	2
Number of Tables	1	2	5
Number of Dishes	1	2	5
Plan Length	5	13	40
BFWS Nodes Generated	11	72	Much higher
BFWS Time (seconds)	0.000144	0.000219	Increased
Solution Consistency	BFWS-Delfi agree	BFWS-Delfi agree	BFWS-Delfi agree

We make the following observations.

- **Complexity Growth:** Plan length increases from 5 to 13 to 40 (roughly exponential with problem size), demonstrating how quickly planning complexity grows
- **Planner Agreement:** Both BFWS and Delfi consistently find the same optimal solutions across all complexity levels, validating the PDDL domain design
- **Multi-Robot Coordination:** The introduction of a second robot creates branching factors that increase search space significantly, but optimal solutions are still found
- **Routing Efficiency:** The planner makes intelligent decisions about which robot serves which table based on kitchen availability and distance constraints

## 6 Conclusion

The Robot Waiter domain proved to be an effective testbed for automated planning. We developed the PDDL encoding and tested it with three scenarios of increasing complexity, from a simple single-delivery case to a multi-robot, multi-kitchen setup.

The main takeaway is that complexity grows fast. A single table required 5 actions, two tables needed 13, and five tables across two robots required 40. This shows how planning problems scale with the number of agents and tasks. Both BFWS and Delfi found the same optimal solutions in all cases, which confirms our domain design is sound.

The exercise also showed how important graph structure is for planning. The shortcuts and alternative paths available in test case 3 gave the planners more options, but the search complexity increased significantly. In practice, this means that better connectivity in a domain doesn't necessarily make planning easier.

Overall, the Robot Waiter problem is a good learning exercise. It demonstrates how to model realistic coordination tasks in PDDL and how modern planners handle multi-agent scenarios with resource constraints.